

Reinforcement Learning: Policy-based Methods

Benjamin Rosman

School of Computer Science and Applied Mathematics
University of the Witwatersrand
South Africa

African Masters of Machine Intelligence (Kigali, Rwanda)
January 13-17 2020



Benjamin Rosman



Further reading...

- Reinforcement Learning: An Introduction (Sutton and Barto)
 - Chapter 13
- CS 294-122 (Sergey Levine)
 - Lectures 4 and 5
- COMPM050/COMPGI13 (David Silver)
 - Lecture 7


A quick recap

Decision maker (agent) exists within an environment

Agent takes **actions** based on the environment **state**

Environment **state** updates

Agent receives feedback as **rewards**




A model for decision making

Markov Decision Process (MDP)

$$M = \langle S, A, T, R, \gamma \rangle$$

- States: encode world configurations
- Actions: choices made by agent



A model for decision making

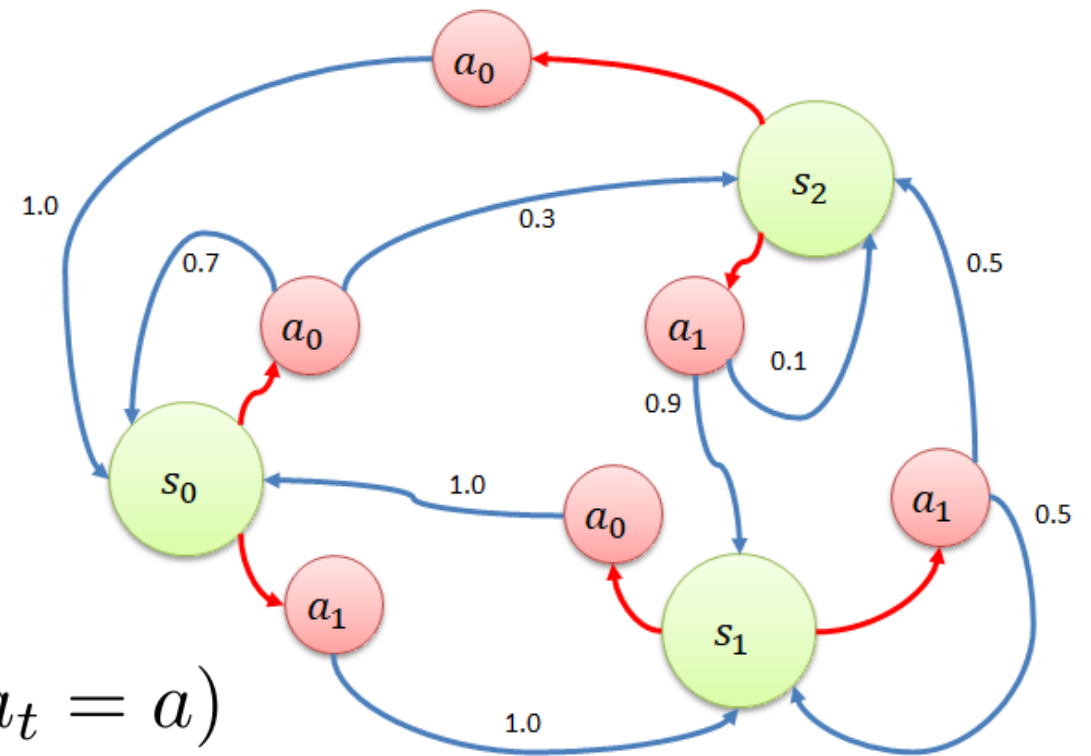
Markov Decision Process (MDP)

$$M = \langle \mathbf{S}, \mathbf{A}, \mathbf{T}, \mathbf{R}, \gamma \rangle$$

- Transition function: how the world evolves under actions

Stochastic!

$$T(s, a, s') = P(s_{t+1} = s' | s_t = s, a_t = a)$$



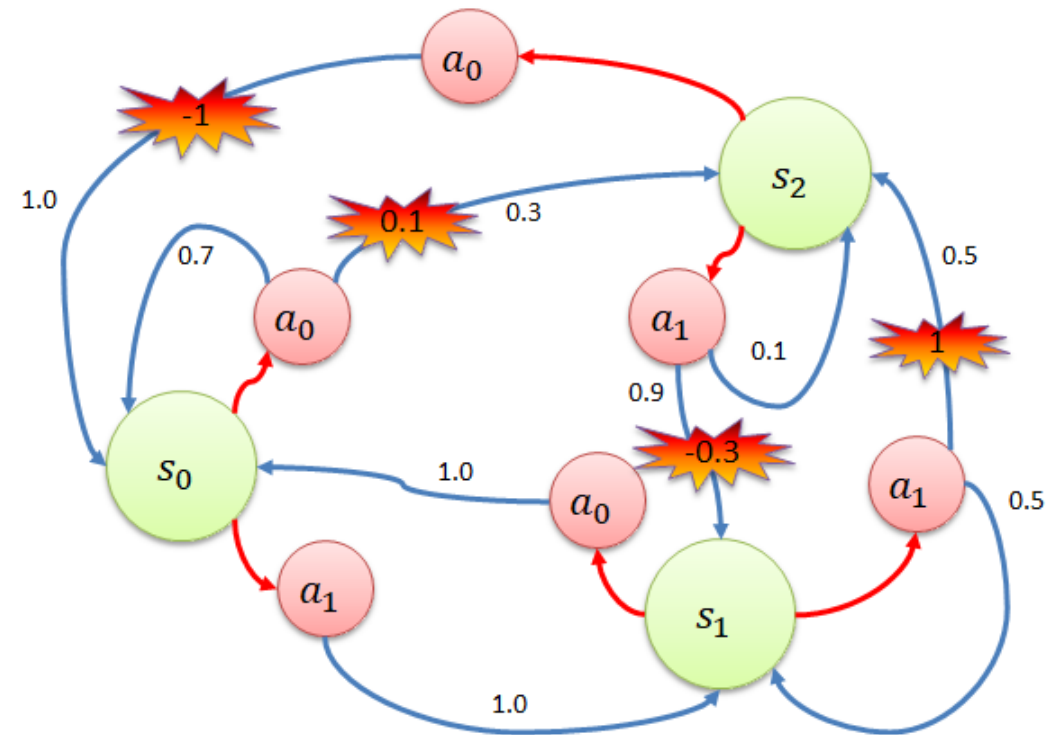
A model for decision making

Markov Decision Process (MDP)

$$M = \langle \mathbf{S}, \mathbf{A}, \mathbf{T}, \mathbf{R}, \gamma \rangle$$

- Rewards: feedback signal to agent

$$R(s, a) = E[r_t | s_t = s, a_t = a]$$



A model for decision making

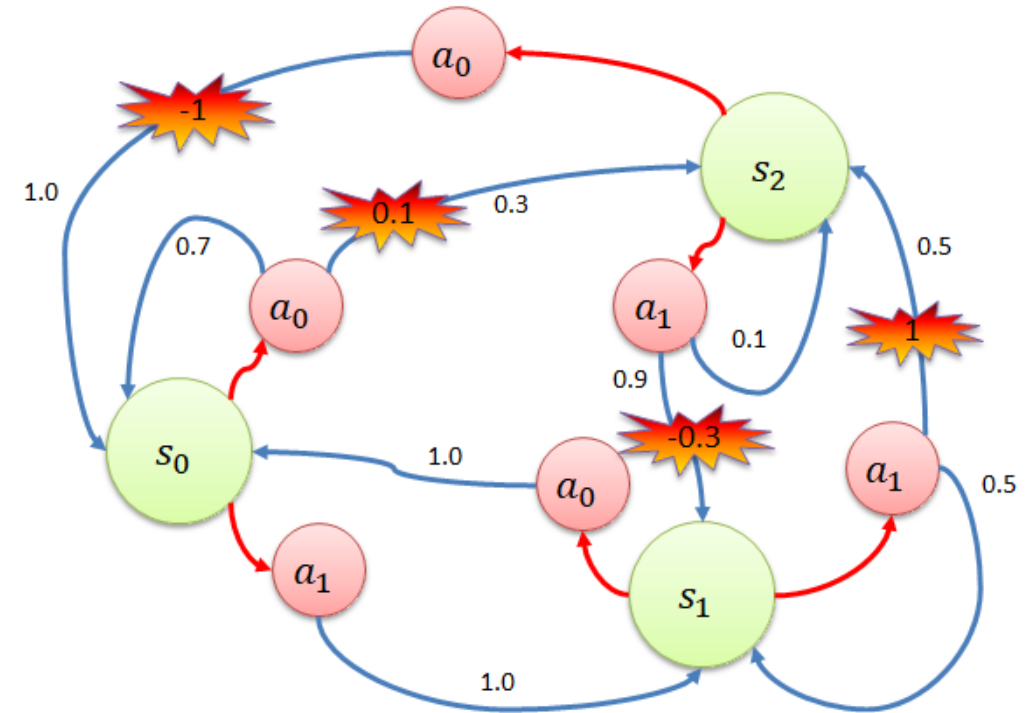
Markov Decision Process (MDP)

$$M = \langle S, A, T, R, \gamma \rangle$$

- $\gamma \in [0,1]$ discounting for future rewards


Particularly important for **continuing** tasks

The value of something now is usually greater than in the future.




Example

- I want to get this robot to walk as far as possible:
 - What is S, A, T, R ?



- I want to play this game as well as possible:
 - What is S, A, T, R ?



Why can't we act greedily?

Cannot just rely on the **instantaneous** reward function

Tradeoff: **don't just act myopically (short term)**



Notion of **value** to codify the goodness of a state, considering a policy running into the future

- Represented as a **value function**


Warning: reward functions

- The reward tells the agent **what** we want it to achieve
 - Learning is about figuring out **how** to achieve it
- Be careful: **the agent will literally do exactly what you ask it**
 - It doesn't "know what you mean"
- Where might this be a problem?

Value functions

Allow us to act greedily by considering estimates of the future


Optimal policy:




a) gridworld

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

b) v_*



c) π_*



NB: Equivalence between optimal policy and optimal value function

Value functions: recursion

$V(s) \Rightarrow$ expected return starting at s and following π
Suggests dependence on value of next state s'

Bellman Equation:

$$\boxed{V^\pi(s)} = \boxed{R(s, \pi(s))} + \gamma \boxed{\sum_{s'}} \boxed{T(s, \pi(s), s')} \boxed{V^\pi(s')}$$

value of s immediate reward for all possible next states the probability of reaching that state with π value of s'

Solution techniques

- If you know T and R :
 - Iteratively compute solution
 - Dynamic programming
 - Policy iteration and value iteration
- If you **don't** know T and R : (more common case)
 - Explore the environment and collect samples
 - Use samples to learn T and R
 - Model-based RL (more next time)
 - Use samples to learn V
 - Model-free RL
 - E.g. TD-learning, SARSA, Q-learning, ...

Q-Learning

- Initialise $Q(s, a)$ for all s, a
- Repeat (for each episode):
 - Initialise s
 - Repeat (for each step of episode):
 - Choose a from s using policy from Q (ϵ -greedy)
 - Take action a , observe r, s'
 - $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
 - $s \leftarrow s'$
- Until s is terminal

ϵ -Greedy ($0 < \epsilon \leq 1$):

- With probability $1 - \epsilon$ exploit:
 - Choose the best action from Q
 - With probability ϵ explore:
 - Randomly choose action
- ϵ usually higher at beginning, decay later

act

learn

take best next action
(so far)

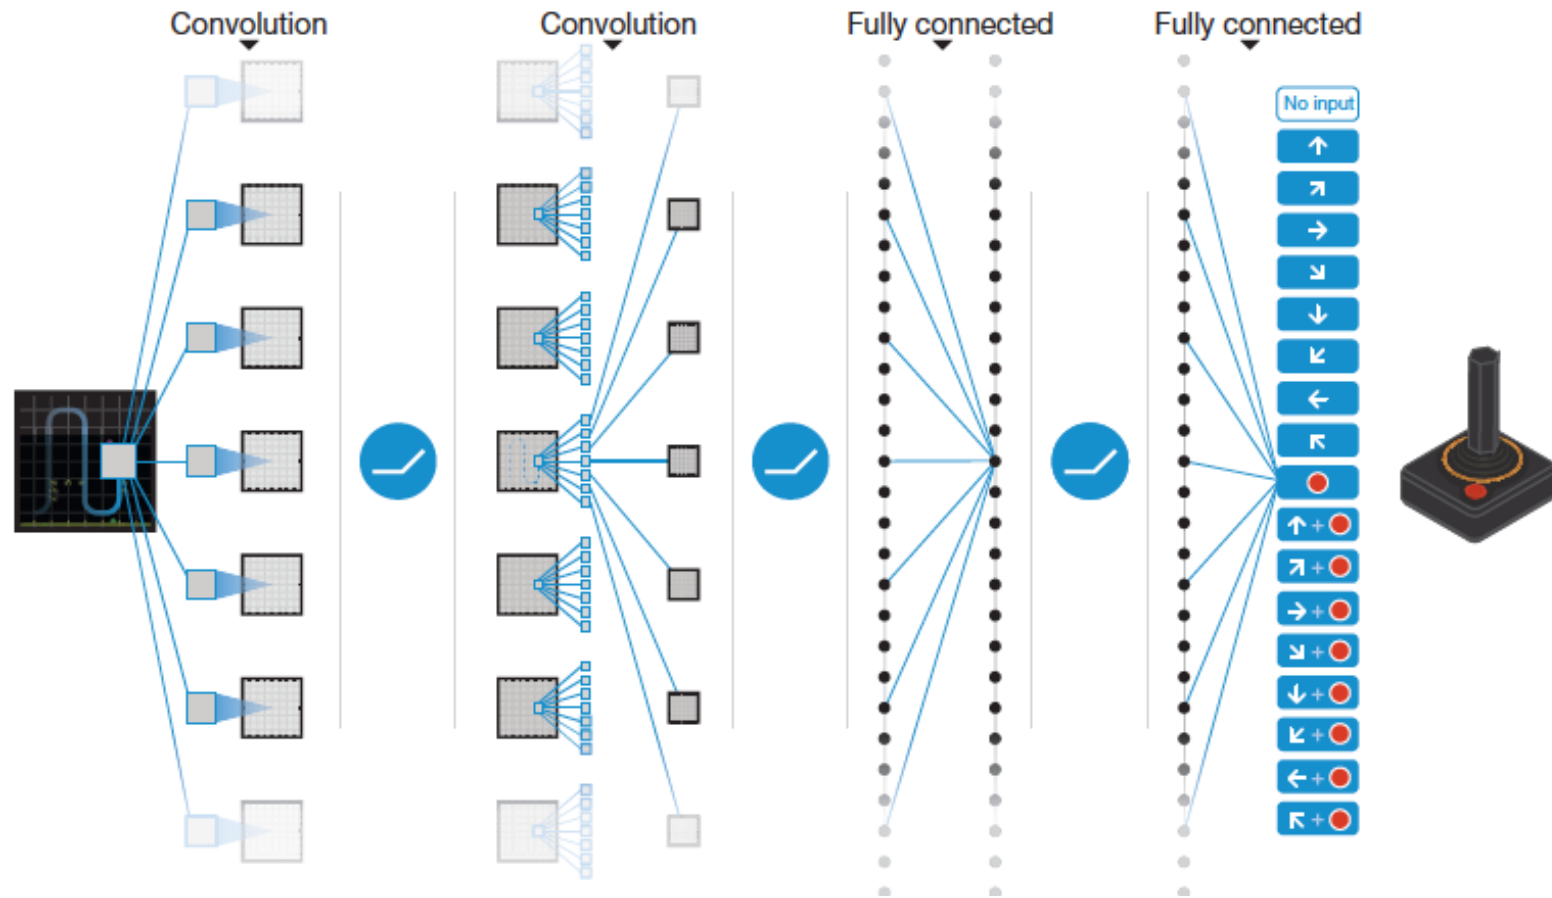
Function approximation

- This works well if you can tabulate the value function

State	a_1	...	a_m
s_1	1.7	...	2.8
s_2	-3.9	...	-3.1
...
s_n	0.2	...	-0.1

- But if the state space is large/continuous
 - Approximate this with a learned function
 - Why?
 - Infinity? May not visit every state (many times)? Need to generalise?

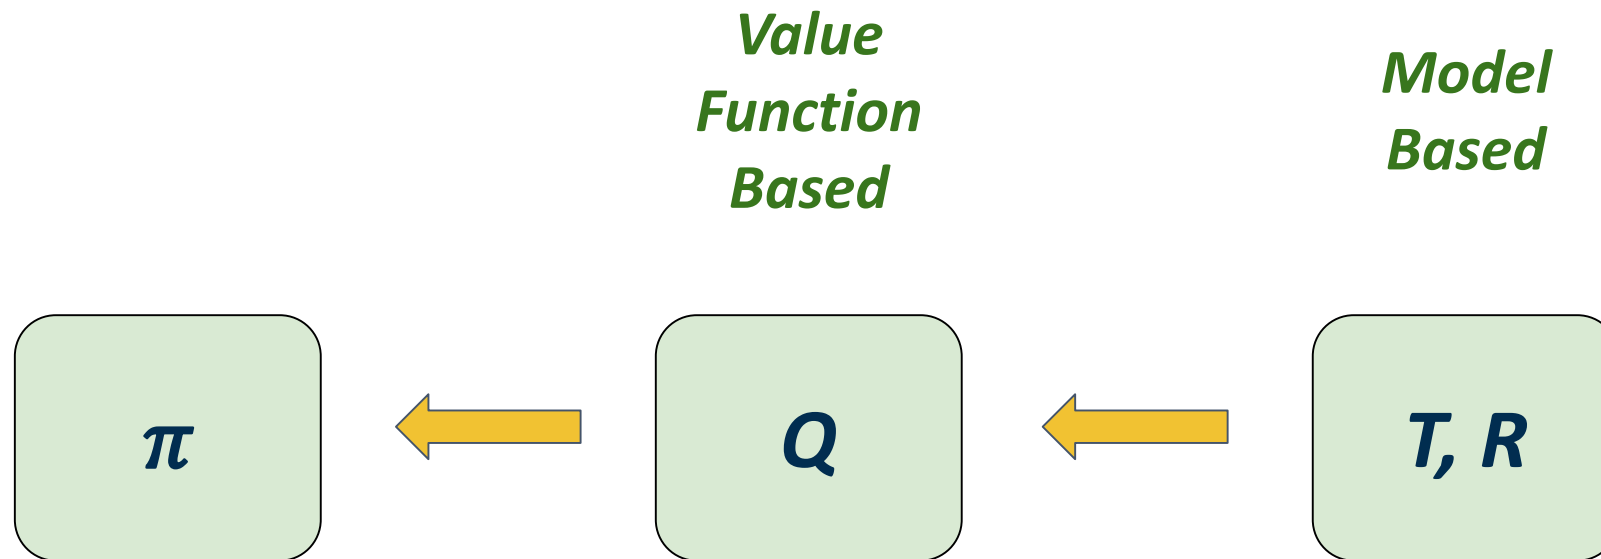
Deep Q-Networks



DQN

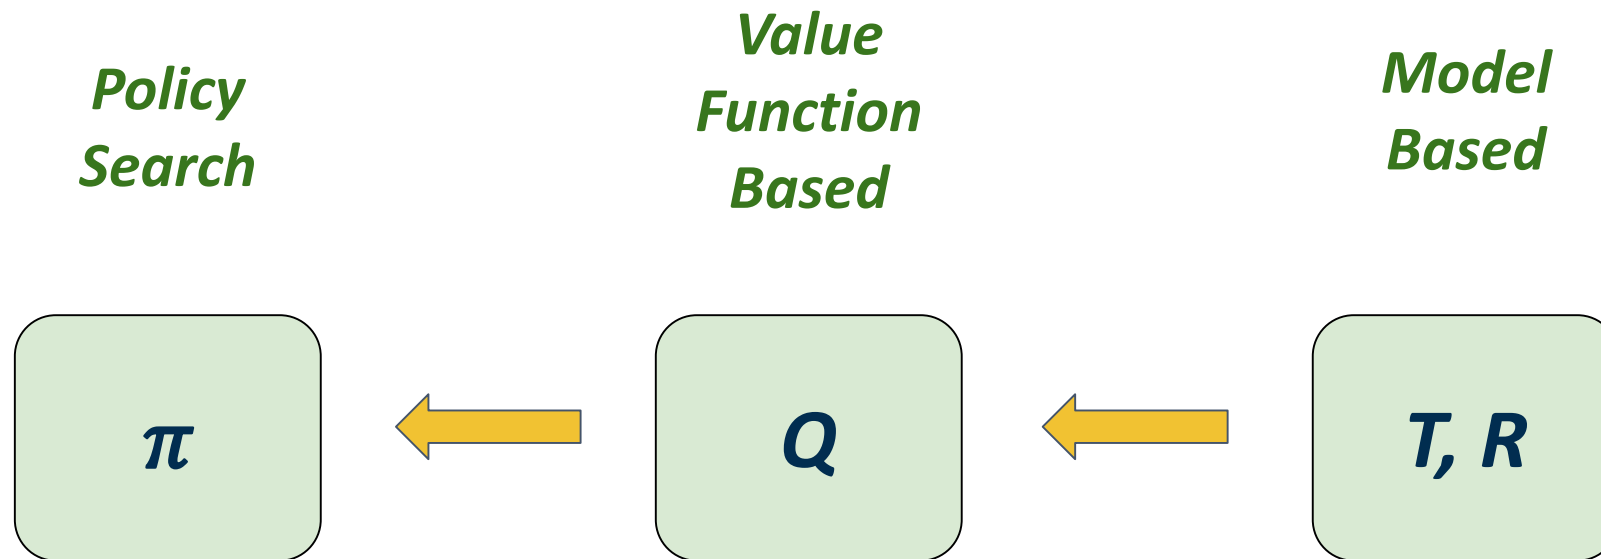
- Take action a_t according to ϵ -greedy policy
 - Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory D
- Sample mini-batch of transitions (s, a, r, s') from D
 - Experience replay: decorrelate samples
- Compute Q-learning targets w.r.t old fixed parameters w^-
 - Fixed Q-targets: avoid oscillations
- Optimise MSE between Q-network and Q-learning targets
 - $L_i(w_i) = \mathbb{E}_{s,a,r,s' \sim D_i} \left[\left(r + \gamma \max_{a'} Q(s', a', w_i^-) - Q(s, a, w_i) \right)^2 \right]$
 - Using stochastic gradient descent on $\nabla_{w_i} L_i(w_i)$

RL approaches



Questions?

RL approaches



Policy Search

- Learn policy directly:

$$\pi_{\theta}(s, a) = \pi(s, a; \theta)$$

- Parameterise policy: learn parameters of policy
- Why?
 - When might it be easier to learn a policy than a value function?
 - Learning a Q-function can be complicated
 - Policy may be much simpler than learning a value for each state-action
 - Injecting information?
 - Stochasticity?

Policy Search

- Objective function?
 - Maximise return given θ :

- $J(\theta) = \mathbb{E}[\sum_t \gamma^t r_t | \pi_\theta]$
 - $\theta^* = \operatorname{argmax}_\theta J(\theta)$

- Always more efficient to follow the gradient!

Trajectory τ

$$\begin{aligned} p(\tau|\theta) &= p_\theta(s_1, a_1, s_2, a_2, \dots, s_T, a_T) \\ &= p(s_1) \prod_{t=1}^T \pi_\theta(a_t | s_t) p(s_{t+1} | s_t, a_t) \end{aligned}$$

Note:

The return R , cost J , utility U are often used interchangeably here

Hill Climbing

- What if you can't differentiate π ?
- Sample-based optimisation:
 - Sample some θ values near your current best θ
 - Compute return
 - Approximate a gradient
 - Finite differences
 - Adjust your current best θ to give the highest value
- Other approaches, e.g. genetic algorithms

Aibo gait optimization

from Kohl and Stone, ICRA 2004




Fig. 2. The elliptical locus of the Aibo's foot. The half-ellipse is defined by length, height, and position in the x - y plane.


All told, the following set of 12 parameters define the Aibo's gait [10]:

- The front locus (3 parameters: height, x -pos., y -pos.)
- The rear locus (3 parameters)
- Locus length
- Locus skew multiplier in the x - y plane (for turning)
- The height of the front of the body
- The height of the rear of the body
- The time each foot takes to move through its locus
- The fraction of time each foot spends on the ground



Using gradients

- If we can differentiate π
 - Compute and ascend $\partial R / \partial \theta$
 - This is the gradient of return w.r.t policy parameters
 - $\theta_{t+1} = \theta_t + \Delta \theta_t$
 $= \theta_t + \alpha \nabla J(\theta_t)$
 - These are called **policy gradient methods**



Policy Gradient Theorem

- Why does this work?
- Relate the **gradient of performance with respect to the policy parameter** to the **gradient of the policy**
- Policy gradient theorem:
 - $\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q^{\pi_{\theta}}(s, a)]$
- No explicit dependence on distribution of states (or model)

What is a good form for a parameterised function f ?

- Simplest thing you can do?
 - *Linear value function approximation*
 - Use set of basis functions ϕ_1, \dots, ϕ_n
 - f is a linear function of them:
 - $\hat{f} = \mathbf{w} \cdot \Phi(s, a) = \sum_{i=1}^n w_i \phi_i(s, a)$
 - We'll want to learn parameters \mathbf{w}
- Neural network:

Basis functions $\phi(x)$:

- Could be polynomial in state vars:
 - 1st order: $[1, x, y]$
 - 2nd order: $[1, x, y, x^2, y^2, xy]$
 - This is a Taylor expansion
- Others:
 - Fourier basis
 - Wavelet basis
 - ...

REward Increment = Nonnegative Factor times Offset
Reinforcement times Characteristic Eligibility



REINFORCE (Monte-Carlo Policy Gradient)

- REINFORCE: one particularly popular sample-based estimate of the gradient
 - Based on the policy gradient theorem, but approximate the \mathbb{E} with sampled trajectories (Monte-Carlo samples)

The return R_t acts as an estimate of $Q^{\pi_\theta}(s, a)$

$$\Delta\theta_t = \alpha R_t \frac{\nabla_\theta \pi_\theta(a_t | s_t)}{\pi_\theta(a_t | s_t)} = \alpha R_t \nabla_\theta \log \pi_\theta(a_t | s_t)$$

Log derivative trick: $\frac{\nabla_\theta x}{x} = \nabla_\theta \log x$

REINFORCE algorithm

- Initialise θ
- For each episode
 - Choose actions according to $\pi_\theta: a \sim \pi_\theta(a|s)$
 - Gather samples $\{s_1, a_1, r_1, \dots, s_T, a_T, r_T\}$
 - For $t = 1$ to T
 - $\theta \leftarrow \theta + \alpha R_t \nabla_\theta \log \pi_\theta(a_t|s_t)$

(that's it)

Deriving REINFORCE

- Cost: $J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau)]$
 $= \int_{\tau} r(\tau) p(\tau; \theta) d\tau$ $\tau = (s_0, a_0, r_0, s_1, \dots)$
- Derivative: $\nabla_{\theta} J(\theta) = \int_{\tau} r(\tau) \nabla_{\theta} p(\tau; \theta) d\tau$
- Transformation: $\nabla_{\theta} p(\tau; \theta) = \boxed{p(\tau; \theta)} \frac{\nabla_{\theta} p(\tau; \theta)}{\boxed{p(\tau; \theta)}} = p(\tau; \theta) \nabla_{\theta} \log p(\tau; \theta)$
Log derivative trick: $\frac{\nabla_{\theta} x}{x} = \nabla_{\theta} \log x$
- Substitute: $\nabla_{\theta} J(\theta) = \int_{\tau} (r(\tau) \nabla_{\theta} \log p(\tau; \theta)) p(\tau; \theta) d\tau$
 $= \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau) \nabla_{\theta} \log p(\tau; \theta)]$


Deriving REINFORCE

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau) \nabla_{\theta} \log p(\tau; \theta)]$$

- Recall: $p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$
- So: $\log p(\tau; \theta) = \sum_{t \geq 0} \log p(s_{t+1} | s_t, a_t) + \log \pi_{\theta}(a_t | s_t)$
- Derivative: $\nabla_{\theta} \log p(\tau; \theta) = \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$ Note we lose dependence on dynamics
- And so estimate: $\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$

Interpretation

- Gradient estimator: $\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$
- Think of this as saying:
 - If $r(\tau)$ is high: push up probabilities of seen actions
 - If $r(\tau)$ is low: push down probabilities of seen actions
- Simple version of credit assignment



Variance

- This gradient estimator (MC) turns out to have a high variance

$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- Why?
 - These are all samples of a run of a policy!
 - Slow convergence
- Correct with a baseline:
- $\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} (r(\tau) - b(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$
- Because b could be 0, this is a generalisation of REINFORCE
 - Will converge asymptotically to a local minimum

Why can we use a baseline?

- $\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} (r(\tau) - b(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$
- Intuition: can add or subtract b from r without biasing algorithm
 - As long as b not a function of a_t

- Mathematically (with some notational abuse):

$$\begin{aligned} & \mathbb{E}_{\pi_{\theta}} \left[\sum_t b \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] \\ &= \int \left[\sum_t b \pi_{\theta}(a_t | s_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] d\tau \end{aligned}$$

$$= \int b \nabla_{\theta} \pi_{\theta}(\tau) d\tau$$

$$= b \nabla_{\theta} \int \pi_{\theta}(\tau) d\tau = b \nabla_{\theta} 1 = 0$$


Keeps the gradient unbiased,
i.e. doesn't change the
expected value, but can
change the variance!

Log derivative trick: $\frac{\nabla_{\theta} x}{x} = \nabla_{\theta} \log x$

Choice of baseline


- What baseline to use?
- $b(s_t) = V(s_t)$
 - Change based on whether or not reward was better than expected
 - Term $r(\tau) - b(s_t)$ resembles advantage function:
 - $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$
 - Measures how much better a is than whatever π would have done
- Suggests we should be learning π and V !

Learning with a baseline



Actor-Critic

- Combine ideas from **policy** and **value function** methods
 - Approximate both the **policy** and the **value function**
- **Actor improvement**
 - Policy parameterised by θ
- **Critic evaluation**
 - Value function parameterised by ω
 - Either $V(s; \omega)$ or $Q(s, a; \omega)$
- Keep track of two sets of parameters



Actor-Critic pseudocode

- Input: parameterised forms for $\pi_{\theta}(s|a)$ and $V_{\omega}(s)$ What forms could we use?

- Input: learning rates $\alpha_{\omega} > 0$ and $\alpha_{\theta} > 0$

- For each episode:

- Initialise s

- For each time step:

- Choose $a \sim \pi_{\theta}(s|a)$

Policy is stochastic: this is a random draw

- Take a , observe s', r

- $\delta \leftarrow r + \gamma V_{\omega}(s') - V_{\omega}(s)$


Compute TD error

- $\omega \leftarrow \omega + \alpha_{\omega} \delta \nabla_{\omega} V_{\omega}(s)$

- $\theta \leftarrow \theta + \alpha_{\theta} \delta \nabla_{\theta} \log \pi_{\theta}(a|s)$

Update parameters by gradient ascent

- $s \leftarrow s'$



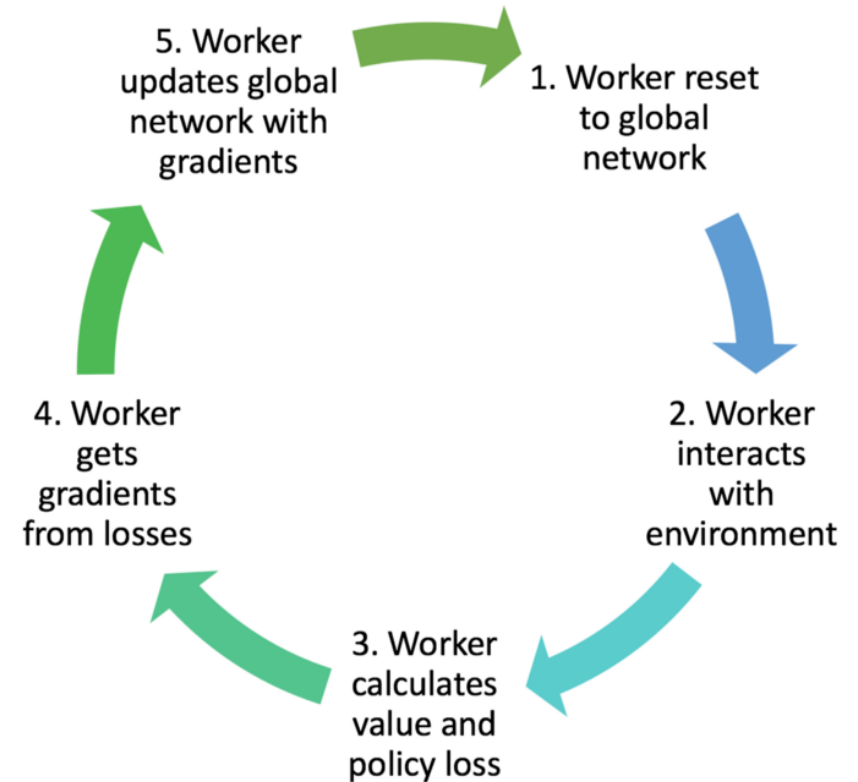
Many different ways to do the updates

$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \textcolor{red}{v}_t]$	REINFORCE
$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \textcolor{red}{Q}^w(s, a)]$	Q Actor-Critic
$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \textcolor{red}{A}^w(s, a)]$	Advantage Actor-Critic
$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \textcolor{red}{\delta}]$	TD Actor-Critic
$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \textcolor{red}{\delta e}]$	TD(λ) Actor-Critic
$G_{\theta}^{-1} \nabla_{\theta} J(\theta) = w$	Natural Actor-Critic

Asynchronous Advantage Actor-Critic (A3C)

[Mnih et al., 2016]

- Actor-Critic can be easily parallelised
- Why is this useful?
 - Speed up exploration of state space
- Have multiple agents training with shared parameters



Algorithm S3 Asynchronous advantage actor-critic - pseudocode for each actor-learner thread.

// Assume global shared parameter vectors θ and θ_v and global shared counter $T = 0$

// Assume thread-specific parameter vectors θ' and θ'_v

Initialize thread step counter $t \leftarrow 1$

repeat

Reset gradients: $d\theta \leftarrow 0$ and $d\theta_v \leftarrow 0$.

Synchronize thread-specific parameters $\theta' = \theta$ and $\theta'_v = \theta_v$

$t_{start} = t$

Get state s_t

repeat

Perform a_t according to policy $\pi(a_t|s_t; \theta')$

Receive reward r_t and new state s_{t+1}

$t \leftarrow t + 1$

$T \leftarrow T + 1$

until terminal s_t **or** $t - t_{start} == t_{max}$

$R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, \theta'_v) & \text{for non-terminal } s_t // \text{ Bootstrap from last state} \end{cases}$

for $i \in \{t - 1, \dots, t_{start}\}$ **do**

$R \leftarrow r_i + \gamma R$

Accumulate gradients wrt θ' : $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$

Accumulate gradients wrt θ'_v : $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

end for

Perform asynchronous update of θ using $d\theta$ and of θ_v using $d\theta_v$.

until $T > T_{max}$

Spin up a new agent/thread

Use global parameters

Act

Update local parameters using
advantage functions

Update global parameters

Deep policy search



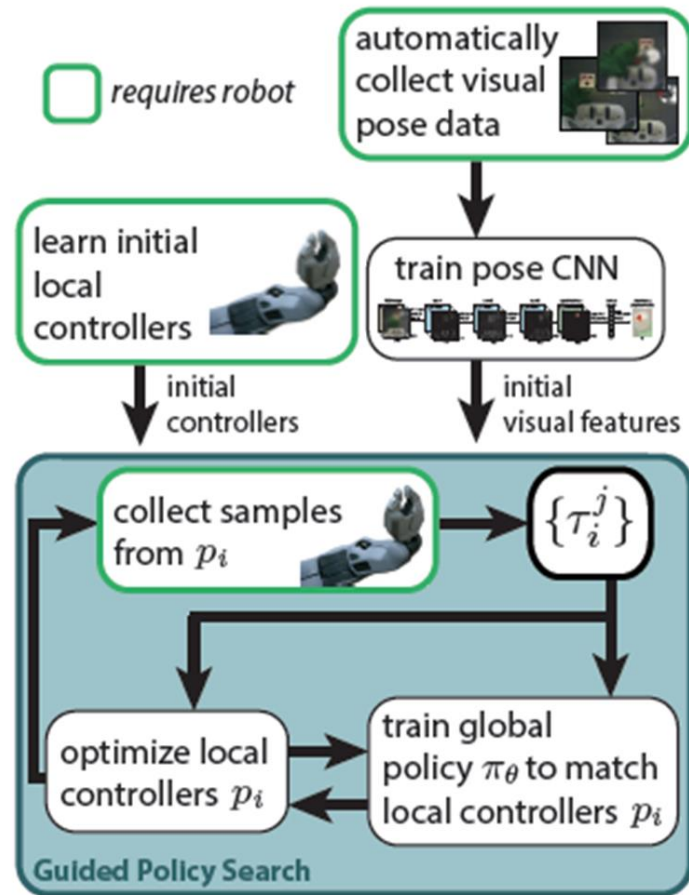


Figure 1: Our method learns visuomotor policies that directly use camera image observations (left) to set motor torques on a PR2 robot (right).



[Levine et al., 2016]

Deep policy search



[Levine et al., 2016]

Robotics

Learned Visuomotor Policy: Shape sorting cube

Conclusion

- Recap:
 - The RL setting, MDPs
 - Rewards and value functions
 - Q-learning
 - Function approximation → DQN
- Policy-based methods
 - Gradient free
 - Hill climbing
 - Gradient based
 - Policy Gradient Theorem
 - REINFORCE
 - Baselines
 - Actor-Critic (A3C)
 - Incorporate ideas from supervised learning, deep learning, etc.